

How to Troubleshoot Log Rotation Issue

One of the problems, when comes to the log rotation, is that the log file does not seem rotated as the configuration file states. It usually ends up with one big log file. The expected log file should be relatively smaller. In this article, I am going to walk you through log rotation.

If we have a production-ready Django application, there are mainly three kinds of log rotation that we would be interested in. One is the application log. The second is the server log. Depending on the server, it could be uWSGI, Gunicorn, etc. In this article, I am going to use uWSGI as an example. The third one is the [logrotate utility](#) which is installed by default on Ubuntu.

Sample Project

In this article, I am going to use a sample project to demonstrate. I uploaded it on my Github. If you are interested, please check out the following Git repository.

<https://github.com/slow999/DjangoAndLogrotate>

Log Rotation 101

How does log rotation work? There are a few ways. The most common way is that when a rotation is triggered, it deletes the oldest copy, renames the rest of the copies along with the current one by the next sequence number, and creates a new log file. For example, the current log file is **mysite.log**. After the first rotation, the list of files would be **mysite.log** and **mysite.log.1**. After the second rotation, the list of files would be **mysite.log**, **mysite.log.1** and **mysite.log.2**. Let's say we make the number of copies being 2. Thus, after the third rotation, the list of files would still be **mysite.log**, **mysite.log.1**, and **mysite.log.2**.

Here we should know [inode](#) a little bit. As the log rotation runs, a new **mysite.log** is created and associated with a new inode. If a file is deleted, its inode won't recycle. Therefore, we can use inode to verify if log rotation is running properly. Next, I am going to walk you through these three types of log rotation using our sample project.

Django/Python Log Rotation

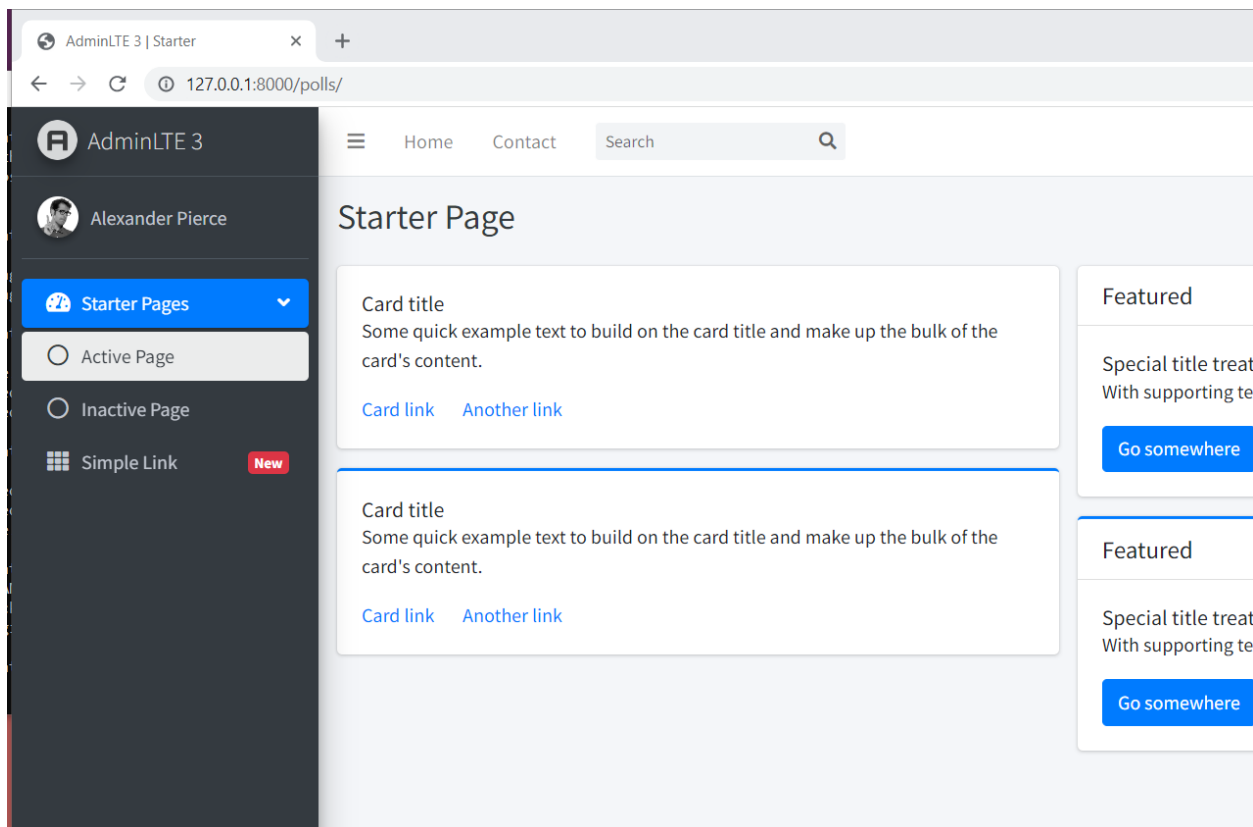
To launch the project, type the command as follows:

```
docker compose up -d
```

We would expect the output in the console as follows:

```
C:\Users\slow9\Documents\MyGit\DjangoAndLogrotate>docker compose up -d
[+] Running 3/3
 - Network djangoandlogrotate_default    Created           0.9s
 - Container djangoandlogrotate-web-1    Started          1.3s
 - Container djangoandlogrotate-nginx-1  Starting         2.4s
```

Open <http://127.0.0.1:8000/polls/> in the browser. We would expect the page as below. It means our project is running.



To display running containers, type the command as follows:

```
docker ps
```

We would expect the output as below. Let's take down the container id that is highlighted in yellow. In my case, it's 0d56a6ed21f0.

```
C:\Users\slow9\Documents\MyGit\DjangoAndLogrotate>docker ps
CONTAINER ID   IMAGE                  COMMAND                  CREATED        STATUS        PORTS                    NAMES
6d2c33b960a2   djangoandlogrotate_nginx  "/docker-entrypoint..." 2 minutes ago  Up 2 minutes  0.0.0.0:8000->80/tcp     djangoandlogrotate-nginx-1
0d56a6ed21f0   djangoandlogrotate_web    "uwsgi --ini mysite..." 2 minutes ago  Up 2 minutes  8000/tcp                 djangoandlogrotate-web-1
```

To get into a running container, type the command as follows:

```
docker exec -it 0d56a6ed21f0 sh
```

We would expect the following output. It means we are in the container.

```
C:\Users\slow9\Documents\MyGit\DjangoAndLogrotate>docker exec -it 0d56a6ed21f0 sh
/code #
```

Let's list what is inside of **code** directory by typing the command as follows:

ls -li

We would expect the following output. Please note that the first column represents the inode. The inode of **mysite.log** is 43132. The **mysite.log** is generated by the Django application.

```
/code # ls -li
total 40
 43079 -rwxr-xr-x  1 root   root         674 Sep 26 04:04 Dockerfile
 43069 -rwxr-xr-x  1 root   root         647 Apr  3 15:46 manage.py
 43114 drwxr-xr-x  1 root   root        4096 Sep 26 04:47 mysite
 43132 -rw-r--r--  1 root   root         919 Sep 29 03:59 mysite.log
 43089 -rwxr-xr-x  1 root   root         233 Sep 26 04:03 mysite.uwsgi.ini
 43133 drwxr-xr-x  1 root   root        4096 Sep 26 04:47 polls
 32999 -rwxr-xr-x  1 root   root         14 Apr  3 15:46 requirements.txt
 42935 drwxr-xr-x  4 root   root        4096 Sep 26 04:47 static
```

The configuration of application logging is set up in the **mysite/settings.py** file. The portion of log rotation is as below. It means if the file size exceeds 1KB, it would invoke log rotation. Django would keep up to 2 copies of log files.

```
logging.config.dictConfig({
    ...
    'handlers': {
        'logfile': {
            'level': 'DEBUG',
            'class': 'logging.handlers.RotatingFileHandler',
            'filename': 'mysite.log',
            'maxBytes': 1024 * 1, # 1KB
            'backupCount': 2,
            'formatter': 'standard',
        },
        'console': {
            'level': 'DEBUG',
            'class': 'logging.StreamHandler',
            'formatter': 'standard',
        },
    },
    ...
})
```

Let's refresh the <http://127.0.0.1:8000/polls/> page one time. List the **code** directory again. We should expect the following outputs. According to the inode, we know that the file whose inode is 43132 has been renamed from **mysite.log** to **mysite.log.1**. A new log file whose inode is 42870 has been created. In addition, please note that the modified time of both log files is the same.

```
/code # ls -li
total 44
 43079 -rwxr-xr-x  1 root  root      674 Sep 26 04:04 Dockerfile
 43069 -rwxr-xr-x  1 root  root      647 Apr  3 15:46 manage.py
 43114 drwxr-xr-x  1 root  root    4096 Sep 26 04:47 mysite
 42870 -rw-r--r--  1 root  root      856 Sep 29 04:13 mysite.log
 43132 -rw-r--r--  1 root  root      982 Sep 29 04:13 mysite.log.1
 43089 -rwxr-xr-x  1 root  root      233 Sep 26 04:03 mysite.uwsgi.ini
 43133 drwxr-xr-x  1 root  root    4096 Sep 26 04:47 polls
 32999 -rwxr-xr-x  1 root  root      14 Apr  3 15:46 requirements.txt
 42935 drwxr-xr-x  4 root  root    4096 Sep 26 04:47 static
```

Let's refresh the <http://127.0.0.1:8000/polls/> page again. Keep watching the changes of the log files. The output of the third refresh is as follows:

```
/code # ls -li
total 48
 43079 -rwxr-xr-x  1 root  root      674 Sep 26 04:04 Dockerfile
 43069 -rwxr-xr-x  1 root  root      647 Apr  3 15:46 manage.py
 43114 drwxr-xr-x  1 root  root    4096 Sep 26 04:47 mysite
 42872 -rw-r--r--  1 root  root      856 Sep 29 04:27 mysite.log
 42870 -rw-r--r--  1 root  root      919 Sep 29 04:27 mysite.log.1
 43132 -rw-r--r--  1 root  root      982 Sep 29 04:13 mysite.log.2
 43089 -rwxr-xr-x  1 root  root      233 Sep 26 04:03 mysite.uwsgi.ini
 43133 drwxr-xr-x  1 root  root    4096 Sep 26 04:47 polls
 32999 -rwxr-xr-x  1 root  root      14 Apr  3 15:46 requirements.txt
 42935 drwxr-xr-x  4 root  root    4096 Sep 26 04:47 static
```

The output of the fourth refresh is as below. The log file whose inode is 43132 is deleted.

```
/code # ls -li
total 48
 43079 -rwxr-xr-x  1 root  root      674 Sep 26 04:04 Dockerfile
 43069 -rwxr-xr-x  1 root  root      647 Apr  3 15:46 manage.py
 43114 drwxr-xr-x  1 root  root    4096 Sep 26 04:47 mysite
 42873 -rw-r--r--  1 root  root      856 Sep 29 04:28 mysite.log
 42872 -rw-r--r--  1 root  root      919 Sep 29 04:28 mysite.log.1
 42870 -rw-r--r--  1 root  root      919 Sep 29 04:27 mysite.log.2
 43089 -rwxr-xr-x  1 root  root      233 Sep 26 04:03 mysite.uwsgi.ini
 43133 drwxr-xr-x  1 root  root    4096 Sep 26 04:47 polls
 32999 -rwxr-xr-x  1 root  root      14 Apr  3 15:46 requirements.txt
 42935 drwxr-xr-x  4 root  root    4096 Sep 26 04:47 static
```

Let's stay in the console. Next, we will dive into uWSGI.

uWSGI Log Rotation

In our sample project, the uWSGI log rotation is set in the **mysite.uwsgi.ini** file. It looks as below. It indicates that if the file size exceeds 2048 bytes, it would rotate.

```
[uwsgi]
```

```
socket = /tmp/uwsgi/mysite.sock
module = mysite.wsgi
master = true
processes = 2
chmod-socket = 666
logger = file:logfile=/tmp/uwsgi.log,maxsize=2048
vacuum = true
```

To list the uWSGI logs, type the following commands.

```
ls -li /tmp/
```

Refresh the web page a few times. A snapshot of the output is as below. Note that the log file whose inode is 42871 has been renamed while the file names of the rest of the copies stay the same. Since there is no limitation on the maximum number of copies, it would keep generating logs till the disk is full. If you know how to set the maximum number of copies, please leave your comment. It would be appreciated it. In this case, we can create a Cron job that cleans up old copies then.

```
/code # ls -li /tmp/
total 300
 27476 -rw----- 1 root    root      285222 Sep  8 00:46 tmpyi4hmf24
 74262 drwxr-xr-x 2 root    root        4096 Sep 29 03:51 uwsgi
 42871 -rw-r----- 1 root    root       1482 Sep 29 05:29 uwsgi.log
 42023 -rw-r----- 1 root    root      2053 Sep 29 03:59 uwsgi.log.1664423949
 42866 -rw-r----- 1 root    root      2073 Sep 29 04:13 uwsgi.log.1664424789
 42869 -rw-r----- 1 root    root      2071 Sep 29 04:28 uwsgi.log.1664425685
/code # ls -li /tmp/
total 304
 27476 -rw----- 1 root    root      285222 Sep  8 00:46 tmpyi4hmf24
 74262 drwxr-xr-x 2 root    root        4096 Sep 29 03:51 uwsgi
 42870 -rw-r----- 1 root    root         550 Sep 29 05:30 uwsgi.log
 42023 -rw-r----- 1 root    root      2053 Sep 29 03:59 uwsgi.log.1664423949
 42866 -rw-r----- 1 root    root      2073 Sep 29 04:13 uwsgi.log.1664424789
 42869 -rw-r----- 1 root    root      2071 Sep 29 04:28 uwsgi.log.1664425685
 42871 -rw-r----- 1 root    root      2050 Sep 29 05:30 uwsgi.log.1664429405
```

Linux logrotate Utility

Logrotate is a system utility that allows automatic rotation, compression, removal, and mailing of log files. A snapshot of the configuration is as below. It means the logrotate utility runs every day, doesn't write an error message if the log file is missing, keeps 14 old log files, compresses the rotated files, and prerotate/postrotate scripts only run once.

```
/var/log/uwsgi/*.log {
    daily
    missingok
    rotate 14
    compress
    sharedscripts
}
```

For the people who have experience with the very old version of uWSGI, the setup above might look familiar to you. We mentioned previously that uWSGI supports file-size-based rotation. Why do we need the logrotate utility to handle the rotation for the uWSGI log file? The reason is that back then, the old uWSGI server does not support file-size-based rotation. Thus, a workaround is to use logrotate as an alternative.

However, the above configuration has a problem. The uWSGI would keep writing to the same log file (same inode) even though the file has been renamed. To solve this issue, we must restart uWSGI after each rotation. As uWSGI restarts, it should generate a new log file (new inode). The revised configuration is as follows:

```
/var/log/uwsgi/*.log {
    daily
    missingok
    rotate 14
    compress
    sharedscripts
    postrotate
        systemctl stop uwsgi.service >/dev/null 2>&1
        systemctl start uwsgi.service >/dev/null 2>&1
    endscript
}
```

Conclusion

In this article, we went through three kinds of log rotation. They are application log rotation, server log rotation, and system log rotation. The discussion on the Linux logrotate utility does not end. Because the revised configuration has a negative impact. The application would have a hiccup during log rotation. If the business requires the app to be 24/7, it might be an issue. I would like to leave it open to you. Please leave comments and let me know your thoughts. Thanks for reading. Stay tuned.